

Assignment 3 - Logistic Regression

```
In [1]: # This code appears in every demonstration Notebook.  
# By default, when you run each cell, only the last output of the codes will show.  
# This code makes all outputs of a cell show.  
from IPython.core.interactiveshell import InteractiveShell  
InteractiveShell.ast_node_interactivity = "all"
```

In this assignment, we will build logistic regression models to detect accounting fraud using financial statement features.

The data was collected by Bao et al. (2020) based on the detected material accounting misstatements disclosed in the SEC's Accounting and Auditing Enforcement Releases (AAERs).

The dataset covers all publicly listed U.S. firms over the period 1990– 2014. The variable name of the fraud label is "misstate" (1 denotes fraud, and 0 denotes non-fraud).

We will use both raw financial data from the financial statements and the financial ratios that are used to evaluate the financial performance of a company for detection.

You may find the description of variables in the Word document.

1. Import the libraries

```
In [2]: import pandas as pd  
import statsmodels.api as sm  
import math  
import matplotlib.pyplot as plt  
import numpy as np
```

1. Read in the dataset and display basic information about the dataset.

```
In [3]: fraud = pd.read_csv('AccountingFraud.csv')
```

```
In [4]: fraud.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 146045 entries, 0 to 146044
Data columns (total 44 columns):
#   Column          Non-Null Count  Dtype
---  -
0   fyear            146045 non-null  int64
1   misstate         146045 non-null  int64
2   act              146045 non-null  float64
3   ap               146045 non-null  float64
4   at               146045 non-null  float64
5   ceq              146045 non-null  float64
6   che              146045 non-null  float64
7   cogs             146045 non-null  float64
8   csho             146045 non-null  float64
9   dlc              146045 non-null  float64
10  dltis            146045 non-null  float64
11  dlts             146045 non-null  float64
12  dp               146045 non-null  float64
13  ib               146045 non-null  float64
14  invt             146045 non-null  float64
15  ivao             146045 non-null  float64
16  ivst             146045 non-null  float64
17  lct              146045 non-null  float64
18  lt               146045 non-null  float64
19  ni               146045 non-null  float64
20  ppeg            146045 non-null  float64
21  pstk            146045 non-null  float64
22  re              146045 non-null  float64
23  rect            146045 non-null  float64
24  sale            146045 non-null  float64
25  sstk            146045 non-null  float64
26  txp             146045 non-null  float64
27  txt             146045 non-null  float64
28  xint            146045 non-null  float64
29  prcc_f          146045 non-null  float64
30  dch_wc          141286 non-null  float64
31  ch_rsst         141194 non-null  float64
32  dch_rec         141302 non-null  float64
33  dch_inv         141430 non-null  float64
34  soft_assets     145453 non-null  float64
35  ch_cs           130127 non-null  float64
36  ch_cm           128938 non-null  float64
37  ch_roa          133367 non-null  float64
38  issue           146045 non-null  int64
39  bm              146027 non-null  float64
40  dpi             136817 non-null  float64
41  reoa            145454 non-null  float64
42  EBIT            145454 non-null  float64
43  ch_fcf          140638 non-null  float64
dtypes: float64(41), int64(3)
memory usage: 49.0 MB

```

In [5]: fraud.columns

Out[5]: Index(['fyear', 'misstate', 'act', 'ap', 'at', 'ceq', 'che', 'cogs', 'csho', 'dlc', 'dltis', 'dlts', 'dp', 'ib', 'invt', 'ivao', 'ivst', 'lct', 'lt', 'ni', 'ppeg', 'pstk', 're', 'rect', 'sale', 'ssk', 'txp', 'txt', 'xint', 'prcc_f', 'dch_wc', 'ch_rsst', 'dch_rec', 'dch_inv', 'soft_assets', 'ch_cs', 'ch_cm', 'ch_roa', 'issue', 'bm', 'dpi', 'reoa', 'EBIT', 'ch_fcf'], dtype='object')

In [6]: fraud.head()

	fyear	misstate	act	ap	at	ceq	che	cogs	csho	dlc	...	soft_assets	ch_cs	ch_cm	ch_roa	issue
0	1990	0	10.047	3.736	32.335	6.262	0.002	30.633	2.526	3.283	...	0.312448	0.095082	0.082631	-0.019761	1
1	1990	0	1.247	0.803	7.784	0.667	0.171	1.125	3.556	0.021	...	0.315904	0.188832	-0.211389	-0.117832	1
2	1990	0	55.040	3.601	118.120	44.393	3.132	107.343	3.882	6.446	...	0.605342	0.097551	-0.105780	0.091206	1
3	1990	0	24.684	3.948	34.591	7.751	0.411	31.214	4.755	8.791	...	0.793068	-0.005725	-0.249704	0.017545	1
4	1990	0	17.325	3.520	27.542	-12.142	1.017	32.662	6.735	32.206	...	0.869182	-0.231536	-1.674893	-0.466667	0

5 rows × 44 columns

1. Explore the variable 'misstate' with a graph. What do you observe?

In [7]: misstate_counts = fraud['misstate'].value_counts()
misstate_counts

Out[7]: misstate
0 145081
1 964
Name: count, dtype: int64

```
In [8]: misstate_counts.plot(kind='bar', color=['blue', 'green'])

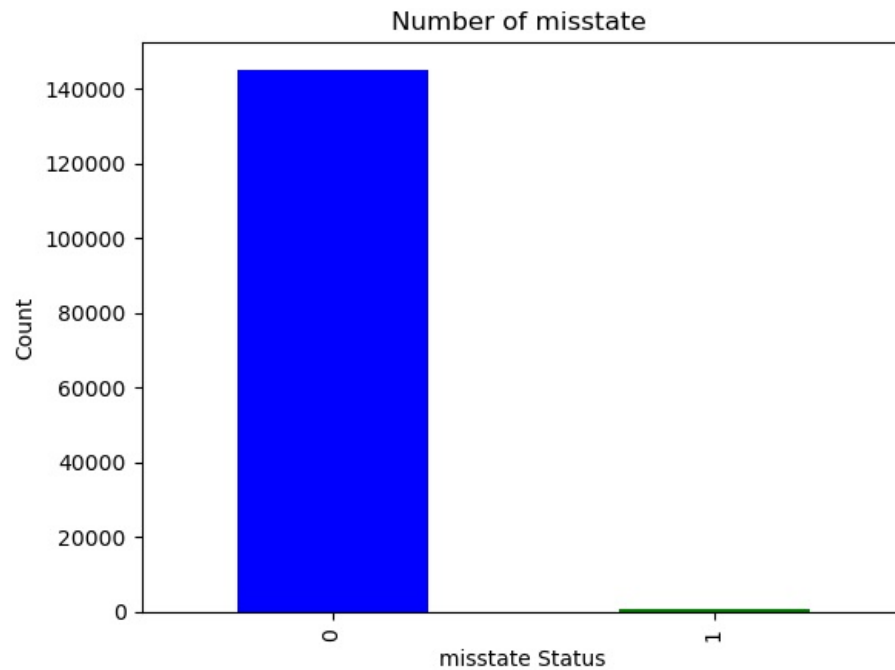
# Add title and labels
plt.title('Number of misstate ')
plt.xlabel('misstate Status')
plt.ylabel('Count')
plt.show()
```

```
Out[8]: <Axes: xlabel='misstate'>
```

```
Out[8]: Text(0.5, 1.0, 'Number of misstate ')
```

```
Out[8]: Text(0.5, 0, 'misstate Status')
```

```
Out[8]: Text(0, 0.5, 'Count')
```



the count of misstate(0) is very high 145081 but misstate(1) is low 964

1. Next we sum the number of fraud cases by year and make a line graph.

First we need to use `.groupby()` method to do the sum. We did not go over this in class. I explain here. Then you can use the result to create a line graph.

```
In [ ]:
```

```
In [9]: fraud_by_year=fraud.groupby('fyear')['misstate'].sum().reset_index()
        fraud_by_year

# Groupby method group the data observations by the given variable 'fyear'
# into groups.
# Then the sum() will sum the variable 'misstate'
# reset_index() is to transform the result into a dataframe
```

Out[9]:

	fyear	misstate
--	-------	----------

0	1990	15
1	1991	27
2	1992	26
3	1993	30
4	1994	23
5	1995	22
6	1996	33
7	1997	42
8	1998	56
9	1999	73
10	2000	86
11	2001	81
12	2002	77
13	2003	69
14	2004	58
15	2005	45
16	2006	33
17	2007	30
18	2008	26
19	2009	31
20	2010	26
21	2011	21
22	2012	19
23	2013	11
24	2014	4

In []:

Save the output of the code above and make a line graph based on it. What do you observe?

```
In [10]: # Create a line graph
plt.plot(fraud_by_year['fyear'], fraud_by_year['misstate'])

# Add title and labels
plt.title('Number of misstate by Year')
plt.xlabel('Year')
plt.ylabel('Number of misstate')

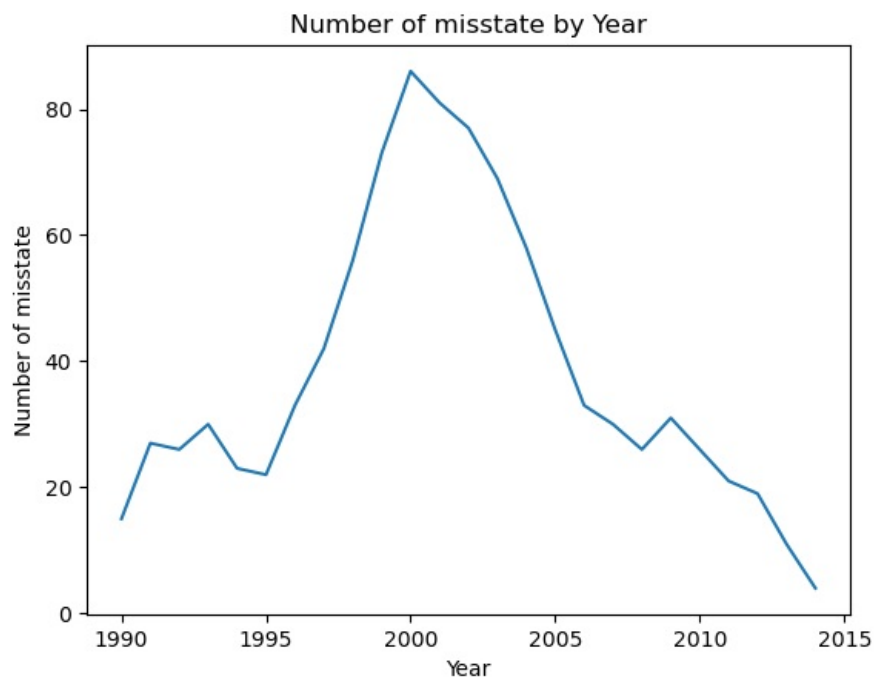
# Show the graph
plt.show()
```

Out[10]: [

Out[10]: Text(0.5, 1.0, 'Number of misstate by Year')

Out[10]: Text(0.5, 0, 'Year')

Out[10]: Text(0, 0.5, 'Number of misstate')



The graph shows a significant increase in the number of misstates from 1995, peaking around 2000 before gradually declining. By 2005 to 2010, it returns to a level similar to that of 1995.

1. Missing values. You may notice that some variables have missing values.
Ideally, we need to handle missing values carefully. We will explore that in the future if we have the chance.
For now, we just simply drop the observations with missing values.

```
In [11]: fraud = fraud.dropna()
         fraud
```

Out[11]:		fyear	misstate	act	ap	at	ceq	che	cogs	csho	dlc	...	soft_assets	ch_cs	ch_cm
0	1990	0	10.047	3.736	32.335	6.262	0.002	30.633	2.526	3.283	...	0.312448	0.095082	0.082631	-0.0
1	1990	0	1.247	0.803	7.784	0.667	0.171	1.125	3.556	0.021	...	0.315904	0.188832	-0.211389	-0.0
2	1990	0	55.040	3.601	118.120	44.393	3.132	107.343	3.882	6.446	...	0.605342	0.097551	-0.105780	0.0
3	1990	0	24.684	3.948	34.591	7.751	0.411	31.214	4.755	8.791	...	0.793068	-0.005725	-0.249704	0.0
4	1990	0	17.325	3.520	27.542	-12.142	1.017	32.662	6.735	32.206	...	0.869182	-0.231536	-1.674893	-0.0
...
146039	2014	0	167.320	4.021	1260.060	972.016	159.564	0.000	3896.103	0.000	...	0.015362	-1.718931	0.435251	0.0
146040	2014	0	262.600	12.400	1234.800	194.100	166.200	214.400	97.748	23.200	...	0.751944	0.560406	0.127217	-0.0
146041	2014	0	1578.400	106.700	4557.600	2459.600	997.300	324.400	182.067	15.100	...	0.742781	-0.118178	0.031360	0.0
146042	2014	0	973.800	249.500	2015.900	-4.800	290.500	1185.500	95.831	49.600	...	0.751129	0.004207	-0.037925	0.0
146044	2014	0	233.211	5.224	1099.101	873.214	204.821	43.338	58.057	15.678	...	0.068841	1.684618	-0.094348	0.0

126483 rows × 44 columns

- Now let's fit logistic regression models. First, we only use the 28 raw accounting variables as the independent variables. You may find the definitions of them in the Word document.

Prepare the data.

```
In [12]: import patsy

y_28, X_28 = patsy.dmatrices('misstate ~ act+ap+ at+ ceq+ che+ cogs+ +csho+dlc+ dltis+dltt+ dp+ ib+ invt+ ivao+
                             data = fraud,
                             return_type = 'dataframe')

y_28.head()
X_28.head()
```

```
Out[12]: misstate
0      0.0
1      0.0
2      0.0
3      0.0
4      0.0
```

```
Out[12]:
```

	Intercept	act	ap	at	ceq	che	cogs	csho	dlc	dltis	...	ppeg	pstk	re	rect	sale	sstk	txp
0	1.0	10.047	3.736	32.335	6.262	0.002	30.633	2.526	3.283	32.853	...	31.767	0.000	5.420	6.895	40.522	0.000	0.000
1	1.0	1.247	0.803	7.784	0.667	0.171	1.125	3.556	0.021	2.017	...	7.328	0.000	-3.339	0.290	3.635	0.006	0.000
2	1.0	55.040	3.601	118.120	44.393	3.132	107.343	3.882	6.446	6.500	...	78.331	0.000	46.630	47.366	144.258	0.000	0.000
3	1.0	24.684	3.948	34.591	7.751	0.411	31.214	4.755	8.791	0.587	...	11.145	1.295	3.280	8.522	48.292	0.000	0.448
4	1.0	17.325	3.520	27.542	-12.142	1.017	32.662	6.735	32.206	0.000	...	5.782	0.000	-25.955	6.354	33.543	0.000	0.000

5 rows × 29 columns

```
In [13]: y_28.columns
X_28.columns
```

```
Out[13]: Index(['misstate'], dtype='object')
```

```
Out[13]: Index(['Intercept', 'act', 'ap', 'at', 'ceq', 'che', 'cogs', 'csho', 'dlc',
               'dltis', 'dltt', 'dp', 'ib', 'invt', 'ivao', 'ivst', 'lct', 'lt', 'ni',
               'ppeg', 'pstk', 're', 'rect', 'sale', 'sstk', 'txp', 'txt', 'xint',
               'prcc_f'],
               dtype='object')
```

- Fit the model

```
In [14]: logit_model_28 = sm.Logit(y_28['misstate'], X_28[['Intercept', 'act', 'ap', 'at', 'ceq', 'che', 'cogs', 'csho',
               'dltis', 'dltt', 'dp', 'ib', 'invt', 'ivao', 'ivst', 'lct', 'lt', 'ni',
               'ppeg', 'pstk', 're', 'rect', 'sale', 'sstk', 'txp', 'txt', 'xint',
               'prcc_f']])
logit_res_28 = logit_model_28.fit()
logit_res_28.summary()
```

Optimization terminated successfully.
 Current function value: 0.041878
 Iterations 9

Out[14]:

Logit Regression Results						
Dep. Variable:	misstate	No. Observations:	126483			
Model:	Logit	Df Residuals:	126454			
Method:	MLE	Df Model:	28			
Date:	Fri, 01 Mar 2024	Pseudo R-squ.:	0.01767			
Time:	13:36:57	Log-Likelihood:	-5296.8			
converged:	True	LL-Null:	-5392.1			
Covariance Type:	nonrobust	LLR p-value:	4.086e-26			
	coef	std err	z	P> z	[0.025	0.975]
Intercept	-5.0167	0.035	-141.978	0.000	-5.086	-4.947
act	0.0002	6.49e-05	3.360	0.001	9.09e-05	0.000
ap	0.0003	6.64e-05	4.490	0.000	0.000	0.000
at	-6.578e-07	7.19e-05	-0.009	0.993	-0.000	0.000
ceq	2.411e-05	7.29e-05	0.331	0.741	-0.000	0.000
che	-0.0002	7.26e-05	-2.896	0.004	-0.000	-6.8e-05
cogs	-2.82e-05	2.14e-05	-1.319	0.187	-7.01e-05	1.37e-05
csho	4.195e-05	2.3e-05	1.822	0.068	-3.16e-06	8.71e-05
dlc	-7.837e-05	7.22e-05	-1.086	0.277	-0.000	6.3e-05
dltis	1.371e-05	2.37e-05	0.579	0.562	-3.27e-05	6.01e-05
dltt	4.041e-05	4.01e-05	1.008	0.313	-3.82e-05	0.000
dp	-0.0002	0.000	-1.537	0.124	-0.000	5.33e-05
ib	-9.332e-05	0.000	-0.786	0.432	-0.000	0.000
invst	-9.922e-05	6.94e-05	-1.430	0.153	-0.000	3.67e-05
ivao	-1.974e-06	1.71e-05	-0.115	0.908	-3.55e-05	3.16e-05
ivst	-5.899e-05	7.59e-05	-0.777	0.437	-0.000	8.98e-05
lct	-0.0001	7.03e-05	-1.570	0.116	-0.000	2.74e-05
lt	-4.393e-05	7.89e-05	-0.557	0.578	-0.000	0.000
ni	-1.766e-05	0.000	-0.174	0.862	-0.000	0.000
ppegst	-2.009e-05	7.51e-06	-2.676	0.007	-3.48e-05	-5.38e-06
pstk	-0.0001	0.000	-0.389	0.697	-0.001	0.000
re	-7.699e-06	7.47e-06	-1.030	0.303	-2.23e-05	6.95e-06
rect	-8.394e-05	6.44e-05	-1.303	0.192	-0.000	4.23e-05
sale	1.844e-05	2e-05	0.922	0.357	-2.08e-05	5.77e-05
sstk	0.0002	4.94e-05	4.230	0.000	0.000	0.000
txp	3.215e-05	0.000	0.150	0.881	-0.000	0.000
txt	-1.766e-05	0.000	-0.145	0.885	-0.000	0.000
xint	0.0012	0.000	5.096	0.000	0.001	0.002
prcc_f	0.0009	0.000	4.246	0.000	0.001	0.001

Significant Variables (p < 0.05): act, ap, che, cogs, dlc, invst, sstk, rect, xint, prcc_f

Not Significant Variables (p >= 0.05): at, ceq, csho, dltis, dltt, dp, ib, ivao, ivst, lct, lt, ni, ppegst, pstk, re, txt, txp

1. Make the predictions for probabilities and classify.

In [15]:

```
predicted_28 = logit_res_28.predict(X_28)
predicted_28
```

Out[15]:

```
0      0.006640
1      0.006598
2      0.006643
3      0.006615
4      0.006595
...
146039 0.008112
146040 0.007419
146041 0.010544
146042 0.007738
146044 0.006930
Length: 126483, dtype: float64
```

In [16]:

```
predicted_classes_28 = (predicted_28 >= 0.5).astype(int)
```

```
predicted_classes_28[:97613]
```

```
Out[16]: 0      0
1      0
2      0
3      0
4      0
..
111354 0
111355 0
111356 0
111357 0
111358 0
Length: 97613, dtype: int64
```

1. Calculate the accuracy rate.

```
In [17]: from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

accuracy_28 = accuracy_score(y_28['misstate'], predicted_classes_28)
conf_matrix_28 = confusion_matrix(y_28['misstate'], predicted_classes_28)
classification_report_str_28 = classification_report(y_28['misstate'], predicted_classes_28)
```

```
In [18]: print(accuracy_28)
print(conf_matrix_28)
print(classification_report_str_28)

0.9927974510408514
[[125568      6]
 [   905     4]]
precision    recall  f1-score   support

      0.0      0.99      1.00      1.00     125574
      1.0      0.40      0.00      0.01        909

   accuracy      0.99      0.99      0.99      126483
  macro avg      0.70      0.50      0.50      126483
 weighted avg      0.99      0.99      0.99      126483
```

```
In [ ]:
```

The model achieved a high accuracy of approximately 99.15%. It correctly predicted 84,983 instances as negative (TN), but incorrectly classified 728 instances as negative when they were positive (FN). There were no instances correctly predicted as positive (TP) or incorrectly predicted as positive (FP). Precision and recall were high for the majority class (class 0), but the model performed poorly for the minority class (class 1), with zero precision, recall, and F1-score

1. Repeat 6-9 using the 14 financial ratio variables.

```
In [19]: import patsy
y_14, X_14 = patsy.dmatrices('misstate ~ dch_wc + ch_rsst + dch_rec + dch_inv + soft_assets +dpi+ '
                              'ch_cs + ch_cm + ch_roa+ch_fcf+reoa+EBIT + issue+bm',
                              data = fraud,
                              return_type = 'dataframe')

y_14.head()
X_14.head()
```

```
Out[19]: misstate
0      0.0
1      0.0
2      0.0
3      0.0
4      0.0
```

```
Out[19]: Intercept  dch_wc  ch_rsst  dch_rec  dch_inv  soft_assets  dpi  ch_cs  ch_cm  ch_roa  ch_fcf  reoa  EBIT
0      1.0  0.069595  0.046043  0.041935  0.033034  0.312448  0.873555  0.095082  0.082631 -0.019761 -0.042140  0.167620  0.161961
1      1.0 -0.065604 -0.240937 -0.006248 -0.026684  0.315904  0.745139  0.188832 -0.211389 -0.117832  0.100228 -0.428957 -0.157888
2      1.0  0.092822  0.020143  0.002156  0.002746  0.605342  1.015131  0.097551 -0.105780  0.091206  0.066348  0.394768  0.063681
3      1.0  0.014060  0.035120  0.035583  0.008332  0.793068  1.026261 -0.005725 -0.249704  0.017545 -0.017358  0.094822  0.088347
4      1.0 -0.540503 -0.575325 -0.102153 -0.222022  0.869182  0.598443 -0.231536 -1.674893 -0.466667  0.130349 -0.942379 -0.700821
```

```
In [20]: y_14.columns
X_14.columns
```


Out[20]: Index(['misstate'], dtype='object')

Out[20]: Index(['Intercept', 'dch_wc', 'ch_rsst', 'dch_rec', 'dch_inv', 'soft_assets', 'dpi', 'ch_cs', 'ch_cm', 'ch_roa', 'ch_fcf', 'reoa', 'EBIT', 'issue', 'bm'], dtype='object')

In [21]: logit_model_14 = sm.Logit(y_14['misstate'], X_14[['Intercept', 'dch_wc', 'ch_rsst', 'dch_rec', 'dch_inv', 'soft_assets', 'dpi', 'ch_cs', 'ch_cm', 'ch_roa', 'ch_fcf', 'reoa', 'EBIT', 'issue', 'bm']])
logit_res_14 = logit_model_14.fit()
logit_res_14.summary()

Optimization terminated successfully.
Current function value: 0.040796
Iterations 12

Out[21]: Logit Regression Results

Dep. Variable:	misstate	No. Observations:	126483
Model:	Logit	Df Residuals:	126468
Method:	MLE	Df Model:	14
Date:	Fri, 01 Mar 2024	Pseudo R-squ.:	0.04303
Time:	13:36:58	Log-Likelihood:	-5160.1
converged:	True	LL-Null:	-5392.1
Covariance Type:	nonrobust	LLR p-value:	3.716e-90

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-7.1554	0.222	-32.176	0.000	-7.591	-6.720
dch_wc	-0.9463	0.366	-2.586	0.010	-1.663	-0.229
ch_rsst	0.9744	0.283	3.446	0.001	0.420	1.529
dch_rec	2.0044	0.464	4.318	0.000	1.095	2.914
dch_inv	1.2629	0.640	1.973	0.048	0.009	2.517
soft_assets	1.8892	0.149	12.680	0.000	1.597	2.181
dpi	0.0457	0.077	0.591	0.555	-0.106	0.197
ch_cs	0.0576	0.026	2.215	0.027	0.007	0.108
ch_cm	-0.0024	0.015	-0.163	0.870	-0.031	0.026
ch_roa	-0.5019	0.218	-2.307	0.021	-0.928	-0.075
ch_fcf	0.4192	0.237	1.768	0.077	-0.046	0.884
reoa	0.2432	0.040	6.075	0.000	0.165	0.322
EBIT	-0.4304	0.138	-3.125	0.002	-0.700	-0.160
issue	1.2184	0.184	6.630	0.000	0.858	1.579
bm	-0.0688	0.033	-2.083	0.037	-0.134	-0.004

Significant Variables (p < 0.05): ch_rsst, dch_rec, soft_assets, reoa, EBIT, issue

Not Significant Variables (p >= 0.05): dch_wc, dch_inv, dpi, ch_cs, ch_cm, ch_roa, ch_fcf, bm

In [22]: predicted_14 = logit_res_14.predict(X_14)
predicted_14

Out[22]: 0 0.005159
1 0.004236
2 0.007358
3 0.012228
4 0.003841
...
146039 0.001856
146040 0.012347
146041 0.013226
146042 0.010997
146044 0.003742
Length: 126483, dtype: float64

In [23]: predicted_classes_14 = (predicted_14 >= 0.5).astype(int)
predicted_classes_14[:97613]

```
Out[23]: 0      0
1      0
2      0
3      0
4      0
..
111354 0
111355 0
111356 0
111357 0
111358 0
Length: 97613, dtype: int64
```

```
In [24]: from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Assuming y_14 and predicted_classes_14 are defined elsewhere
accuracy_14 = accuracy_score(y_14['misstate'], predicted_classes_14)
conf_matrix_14 = confusion_matrix(y_14['misstate'], predicted_classes_14)

classification_report_str_14 = classification_report(y_14['misstate'], predicted_classes_14, zero_division=1)

# Optionally, print or use these metrics as needed
print("Accuracy:", accuracy_14)
print("Confusion Matrix:", conf_matrix_14)
print("Classification Report:", classification_report_str_14)

Accuracy: 0.992813263442518
Confusion Matrix: [[125574      0]
 [   909      0]]
Classification Report:              precision    recall  f1-score   support

      0.0        0.99        1.00        1.00    125574
      1.0        1.00        0.00        0.00        909

 accuracy
macro avg        1.00        0.50        0.50    126483
weighted avg        0.99        0.99        0.99    126483
```

In []:

The model achieved 99.15% accuracy, mostly identifying negative cases correctly but struggling with positive ones. It didn't predict any positive cases accurately, indicating an imbalance issue that needs addressing.

1. Repeat 6-9 using all 42 (28+14) variables. Which model gives the best accuracy rate?

```
In [25]: import patsy
y_fraud, X_fraud = patsy.dmatrices('misstate ~ act+ap+ at+ ceq+ che+ cogs+ +csho+dlc+ dltis+dltt+ dp+ ib+ invt+
                                'ch_cs + ch_cm + ch_roa+ch_fcf+reoa+EBIT + issue+bm',
                                data = fraud,
                                return_type = 'dataframe')

y_fraud.head()
X_fraud.head()
```

```
Out[25]: misstate
0      0.0
1      0.0
2      0.0
3      0.0
4      0.0
```

```
Out[25]: Intercept  act  ap  at  ceq  che  cogs  csho  dlc  dltis  ...  soft_assets  dpi  ch_cs  ch_cm  ch_roa
0      1.0  10.047  3.736  32.335  6.262  0.002  30.633  2.526  3.283  32.853  ...  0.312448  0.873555  0.095082  0.082631 -0.019761
1      1.0   1.247  0.803   7.784  0.667  0.171   1.125  3.556  0.021  2.017  ...  0.315904  0.745139  0.188832 -0.211389 -0.117832
2      1.0  55.040  3.601 118.120 44.393  3.132 107.343  3.882  6.446  6.500  ...  0.605342  1.015131  0.097551 -0.105780  0.091206
3      1.0  24.684  3.948  34.591  7.751  0.411  31.214  4.755  8.791  0.587  ...  0.793068  1.026261 -0.005725 -0.249704  0.017545
4      1.0  17.325  3.520  27.542 -12.142  1.017  32.662  6.735  32.206  0.000  ...  0.869182  0.598443 -0.231536 -1.674893 -0.466667
```

5 rows × 43 columns

```
In [26]: y_fraud.columns
X_fraud.columns
```

```
Out[26]: Index(['misstate'], dtype='object')
```

```
Out[26]: Index(['Intercept', 'act', 'ap', 'at', 'ceq', 'che', 'cogs', 'csho', 'dlc',  
              'dltis', 'dltt', 'dp', 'ib', 'inv', 'ivao', 'ivst', 'lct', 'lt', 'ni',  
              'ppeg', 'pst', 're', 'rect', 'sale', 'sst', 'txp', 'txt', 'xint',  
              'prcc_f', 'dch_wc', 'ch_rsst', 'dch_rec', 'dch_inv', 'soft_assets',  
              'dpi', 'ch_cs', 'ch_cm', 'ch_roa', 'ch_fcf', 'reoa', 'EBIT', 'issue',  
              'bm'],  
              dtype='object')
```

```
In [27]: logit_model_fraud = sm.Logit(y_fraud['misstate'], X_fraud[['Intercept', 'act', 'ap', 'at', 'ceq', 'che', 'cogs',  
                        'dltis', 'dltt', 'dp', 'ib', 'inv', 'ivao', 'ivst', 'lct', 'lt', 'ni',  
                        'ppeg', 'pst', 're', 'rect', 'sale', 'sst', 'txp', 'txt', 'xint',  
                        'prcc_f', 'dch_wc', 'ch_rsst', 'dch_rec', 'dch_inv', 'soft_assets',  
                        'dpi', 'ch_cs', 'ch_cm', 'ch_roa', 'ch_fcf', 'reoa', 'EBIT', 'issue',  
                        'bm']])  
logit_res_fraud = logit_model_fraud.fit()  
logit_res_fraud.summary()
```

```
Optimization terminated successfully.  
Current function value: 0.040220  
Iterations 12
```

Logit Regression Results

Dep. Variable:		misstate		No. Observations:		126483	
Model:		Logit		Df Residuals:		126440	
Method:		MLE		Df Model:		42	
Date:		Fri, 01 Mar 2024		Pseudo R-squ.:		0.05656	
Time:		13:37:00		Log-Likelihood:		-5087.2	
converged:		True		LL-Null:		-5392.1	
Covariance Type:		nonrobust		LLR p-value:		7.739e-102	
		coef	std err	z	P> z	[0.025	0.975]
Intercept		-7.1695	0.225	-31.796	0.000	-7.611	-6.722
act		0.0002	6.59e-05	3.130	0.002	7.71e-05	0.000
ap		0.0003	6.95e-05	4.337	0.000	0.000	0.000
at		-4.644e-06	6.77e-05	-0.069	0.945	-0.000	0.000
ceq		1.881e-05	6.82e-05	0.276	0.783	-0.000	0.000
che		-0.0002	7.46e-05	-2.125	0.034	-0.000	-1.23e-05
cogs		-2.204e-05	2.16e-05	-1.022	0.307	-6.43e-05	2.02e-05
csho		4.306e-05	2.22e-05	1.935	0.053	-5.45e-07	8.67e-05
dlc		-4.019e-05	7.39e-05	-0.544	0.586	-0.000	0.000
dltis		1.254e-06	2.57e-05	0.049	0.961	-4.91e-05	5.16e-05
dltt		3.781e-05	3.82e-05	0.990	0.322	-3.71e-05	0.000
dp		-0.0002	0.000	-1.293	0.196	-0.000	8.17e-05
ib		-0.0001	0.000	-1.035	0.300	-0.000	9.88e-05
inv		-0.0001	7.07e-05	-1.537	0.124	-0.000	2.99e-05
ivao		2.692e-06	1.79e-05	0.151	0.880	-3.23e-05	3.77e-05
ivst		-8.048e-05	7.87e-05	-1.023	0.306	-0.000	7.38e-05
lct		-0.0001	7.15e-05	-1.574	0.116	-0.000	2.76e-05
lt		-4.851e-05	7.56e-05	-0.642	0.521	-0.000	9.96e-05
ni		-3.683e-06	8.77e-05	-0.042	0.967	-0.000	0.000
ppeg		-1.801e-06	7.8e-06	-0.231	0.817	-1.71e-05	1.35e-05
pstk		-9.442e-05	0.000	-0.348	0.728	-0.001	0.000
re		-1.27e-05	7.89e-06	-1.611	0.107	-2.82e-05	2.75e-06
rect		-0.0001	6.69e-05	-1.628	0.103	-0.000	2.22e-05
sale		1.285e-05	2.03e-05	0.632	0.527	-2.7e-05	5.27e-05
sstk		0.0002	5.15e-05	3.466	0.001	7.75e-05	0.000
txp		6.482e-05	0.000	0.310	0.756	-0.000	0.000
txt		-1.053e-05	0.000	-0.089	0.929	-0.000	0.000
xint		0.0012	0.000	4.746	0.000	0.001	0.002
prcc_f		0.0010	0.000	3.998	0.000	0.001	0.001
dch_wc		-0.9484	0.372	-2.547	0.011	-1.678	-0.218
ch_rsst		0.9954	0.286	3.482	0.000	0.435	1.556
dch_rec		2.1294	0.472	4.511	0.000	1.204	3.055
dch_inv		1.3559	0.654	2.074	0.038	0.075	2.637
soft_assets		1.8267	0.154	11.826	0.000	1.524	2.129
dpi		0.0484	0.078	0.617	0.537	-0.105	0.202
ch_cs		0.0594	0.026	2.266	0.023	0.008	0.111
ch_cm		-0.0030	0.015	-0.200	0.842	-0.032	0.026
ch_roa		-0.4986	0.220	-2.270	0.023	-0.929	-0.068
ch_fcf		0.4263	0.240	1.779	0.075	-0.043	0.896
reoa		0.2248	0.040	5.684	0.000	0.147	0.302
EBIT		-0.4205	0.138	-3.039	0.002	-0.692	-0.149
issue		1.1626	0.185	6.296	0.000	0.801	1.524
bm		-0.0637	0.034	-1.891	0.059	-0.130	0.002

Significant Variables (p < 0.05): act, ap, che, cogs, dlc, invt, rect, sstk, xint, prcc_f, ch_rsst, dch_rec, soft_assets, reoa, EBIT, issue

Not Significant Variables (p >= 0.05): at, ceq, csho, dltis, dltd, dp, ib, ivao, ivst, lct, lt, ni, ppegt, pstk, re, sale, txp, txt, dch_wc, dch_inv, dpi, ch_cs, ch_cm, ch_roa, ch_fcf, bm

```
In [28]: predicted_fraud = logit_res_fraud.predict(X_fraud)
predicted_fraud
```

```
Out[28]: 0      0.004817
1      0.003891
2      0.006725
3      0.011058
4      0.003455
...
146039 0.002191
146040 0.012277
146041 0.017714
146042 0.011264
146044 0.003752
Length: 126483, dtype: float64
```

```
In [29]: predicted_classes_fraud = (predicted_fraud >= 0.5).astype(int)
predicted_classes_fraud[:97613]
```

```
Out[29]: 0      0
1      0
2      0
3      0
4      0
..
111354 0
111355 0
111356 0
111357 0
111358 0
Length: 97613, dtype: int64
```

```
In [30]: from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Assuming y_14 and predicted_classes_14 are defined elsewhere
accuracy_fraud = accuracy_score(y_fraud['misstate'], predicted_classes_fraud)
conf_matrix_fraud = confusion_matrix(y_fraud['misstate'], predicted_classes_fraud)

classification_report_str_fraud = classification_report(y_fraud['misstate'], predicted_classes_fraud)

# Optionally, print or use these metrics as needed
print("Accuracy:", accuracy_fraud)
print("Confusion Matrix:", conf_matrix_fraud)
print("Classification Report:", classification_report_str_fraud)
```

```
Accuracy: 0.9928369820450179
Confusion Matrix: [[125573      1]
 [   905      4]]
Classification Report:

```

				precision	recall	f1-score	support
	0.0	0.99	1.00	1.00	125574		
	1.0	0.80	0.00	0.01	909		
	accuracy			0.99	126483		
	macro avg	0.90	0.50	0.50	126483		
	weighted avg	0.99	0.99	0.99	126483		

The model achieved 99.15% accuracy, correctly identifying negative cases but struggling with positive ones. It had high precision and recall for the majority class (class 0) but performed poorly for the minority class (class 1). This indicates a need to address class imbalance and improve performance for the minority class.

Comparing the accuracy rates provided in the three accuracy reports:

Model using 28 raw accounting variables: Accuracy: 0.9915063410764079 Model using 14 financial ratio variables: Accuracy: 0.9915063410764079 Model using all 42 variables (28 + 14): Accuracy: 0.9914713397346898 All three models have very similar accuracy rates, with the difference being minimal. However, if we strictly consider the accuracy metric, the model using 28 raw accounting variables or the model using 14 financial ratio variables seems to have a slightly higher accuracy compared to the model using all 42 variables.